

- =====
- Virtual memory is a technique that allows the execution of processes that may not be completely in memory.
 - Implemented via
 - Demand Paging
 - Demand Segmentation

// ADVANTAGES OF VIRTUAL MEMORY

=====

- Copy-on-Write
 - Copy-on-Write (COW) allows both parent and child processes to initially share the same pages in memory.
 - Only when the page is modified, then child creates a new copy
- Memory Mapped Files
 - A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses

// DEMAND PAGING

=====

- For demand paging, we use a lazy swapper.
 - Bring a page into memory only when it is needed.
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- Mechanism
 - If Page is needed =====> reference to it
 - If invalid reference =====> abort
 - If not-in-memory =====> bring to memory
 - With each page table entry a valid–invalid bit is associated
 - (1 in-memory, 0 page fault)
 - Get an empty frame in Main Memory.
 - Swap page into frame.
 - valid bit = 1.
 - Restart instruction
 - What happens if there is no free frame in Main Memory?

- find some used frame in Main Memory, but not really in use, swap it out(victim frame)(page replacement)
 - Note: Dirty bit
 - Use modify (dirty) bit to reduce overhead of page transfers – only modified pages are written to disk.
- Overhead of page fault
 - Effective Access Time (EAT)
 - $EAT = (1 - p) \times \text{memory access} + p \times (\text{page fault overhead})$
- Page Replacement Algorithms
 - FIFO
 - OPTIMAL
 - LRU

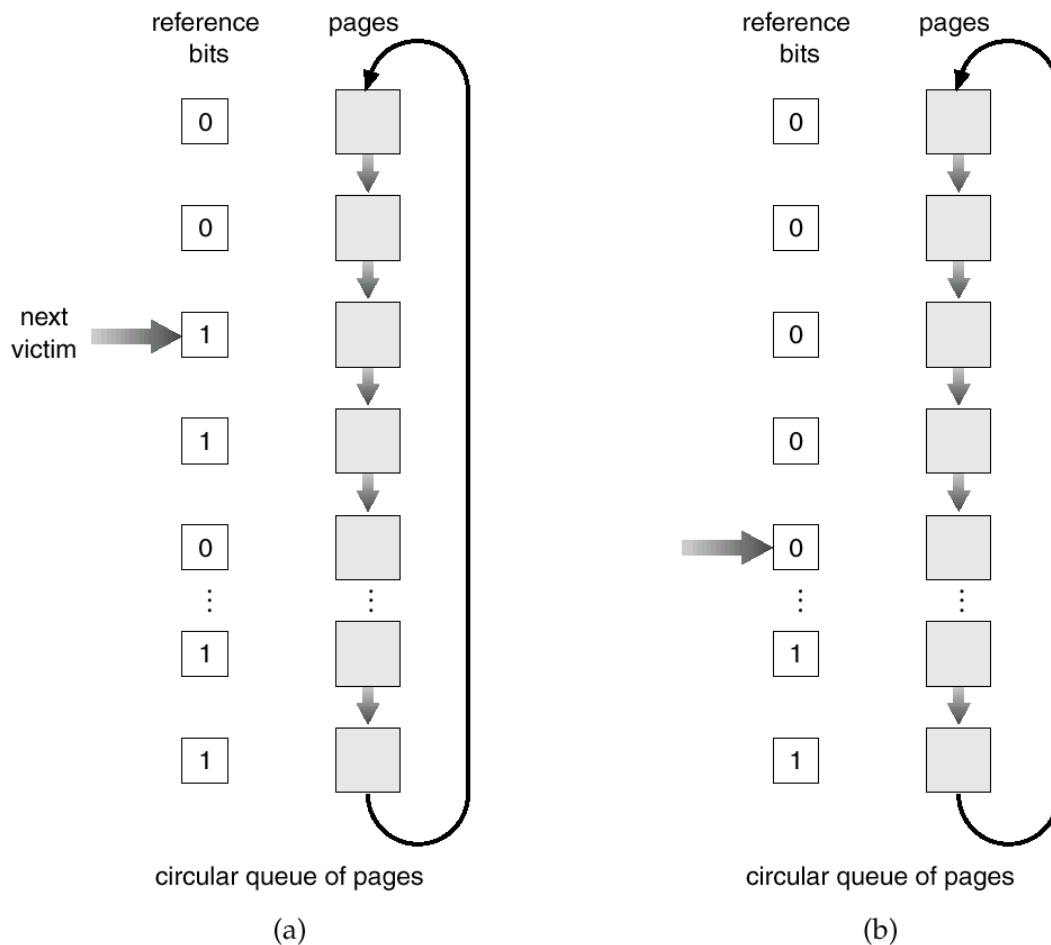
○

// LRU IMPLEMENTATION

=====

- Counter implementation
 - Every page entry has a counter; every time a page is referenced through this entry, copy the clock into the counter.
 - When a page needs to be changed, look at the counters to determine which are to change.
- Stack implementation – keep a stack of page numbers:
 - IF a Page is referenced:
 - move it to the top
 - Look at the stack if a page needs to be changed
- Issues
 - The updating of stack or clock must be done on every memory reference.
 - If we use interrupt for every reference, to allow software to update data structures, it would slow every reference by a factor of 10
- Solutions
 - Reference bit
 - With each page associate a bit, initially = 0
 - When the page is referenced, bit is set to 1.
 - Replace the one which is 0 (if one exists). We do not know the order, however.
 - Second Chance Algorithm

- **Page Reference Bit:**
 - Each page in memory is associated with a reference bit, which is initially set to 0.
- **Page Access:**
 - When a page is accessed (read or written), its reference bit is set to 1, indicating that the page has been used recently.
- **Page Replacement:**
 - When a page needs to be replaced, the algorithm starts scanning the pages in a circular manner. It looks for a page with a reference bit of 0 (not recently used).
 - If it finds a page with a reference bit of 0, it replaces that page.
 - If it finds a page with a reference bit of 1, it gives that page a "second chance" by setting its reference bit to 0 and continues scanning.



//FRAME ALLOCATION SCHEMES - CONTEXT 1

=====

- fixed allocation
 - Equal allocation – e.g., if 100 frames and 5 processes, give each 20 pages.
 - Proportional allocation – Allocate according to the size of the process.
- priority allocation
 - Use a proportional allocation scheme using priorities rather than size.

//FRAME ALLOCATION SCHEMES - CONTEXT 2

=====

- Global replacement – process selects a replacement frame from the set of all frames; one process can take a frame from another.
- Local replacement – each process selects from only its own set of allocated frames

//THRASHING

=====

- Thrashing a process is spending more time swapping pages in and out.
 - If the process does not have # of frames equivalent to # of active pages, it will very quickly page fault.
 - Since all the pages are in active use it will page fault again

//WORKING SET MODEL

=====

- Working Set Window:
 - The working set model defines a "working set window" that represents a period of time. It's a sliding time interval during which the system monitors the pages accessed by a process.
- Working Set Size:
 - The working set size of a process is the number of distinct pages referenced within the working set window. It reflects the temporal locality of the process.
- Working Set Policy:
 - The working set policy involves adjusting the amount of memory allocated to a process based on its working set size. If the working

set size exceeds a certain threshold, the process may be given more memory to improve performance.

//LOCALITY MODEL

=====

Definition:

- The locality model is based on the observation that programs tend to access a small, localized portion of their address space at any given time. This observation is captured by the principle of locality.

Principle of Locality:

Temporal Locality:

- The principle of temporal locality states that if a program accesses a memory location, it is likely to access the same location again in the near future.

Spatial Locality:

- The principle of spatial locality states that if a program accesses a memory location, it is likely to access nearby memory locations in the near future.